

# Notebook

May 12, 2024

Necessary imports

```
[5]: import warnings
import itertools
#from pandas import datetime
#from pandas import read_csv
import numpy as np
import pandas as pd
from pandas.plotting import autocorrelation_plot
from sklearn.metrics import mean_squared_error
from math import sqrt
import statsmodels.api as smf
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
import seaborn as sns
sns.set(style="whitegrid")
```

```
[6]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[7]: import pandas as pd

df_train =df=pd.read_excel("/content/drive/MyDrive/AP moller/DS_ML Coding_
↳Challenge Dataset (1).xlsx",sheet_name='Training Dataset')
```

```
[8]: df_train
```

```
[8]:
```

	ProductType	Manufacturer	Area	Code	Sourcing Channel	Product Size	\
0	NTM3	X1	A28	WHOLESALE	Large		
1	NTM2	X1	A9	DIRECT	Large		
2	NTM3	X2	A20	DIRECT	Large		
3	NTM3	X1	A18	WHOLESALE	Small		
4	NTM2	X1	A28	DIRECT	Large		
...	...	...	...	...	...		
550171	NTM2	X1	A5	DIRECT	Large		
550172	NTM3	X1	A14	DIRECT	Large		
550173	NTM2	X1	A5	DIRECT	Small		
550174	NTM2	X1	A7	DIRECT	Small		
550175	NTM1	X1	A3	DIRECT	Small		

	Product Type	Month of Sourcing	Sourcing Cost
0	Powder	2021-05-01	10.158
1	Powder	2020-10-01	134.281
2	Powder	2020-12-01	12.456
3	Powder	2021-02-01	107.220
4	Liquid	2020-11-01	197.763
...	...	...	...
550171	Powder	2020-07-01	136.469
550172	Liquid	2020-10-01	72.559
550173	Powder	2021-03-01	147.639
550174	Powder	2021-02-01	150.044
550175	Powder	2020-11-01	139.421

[550176 rows x 8 columns]

```
[9]: df_train= df_train.sort_values(by='Month of Sourcing', ascending=False)
```

```
[10]: df_train
```

```
[10]:
```

	ProductType	Manufacturer	Area	Code	Sourcing Channel	Product Size	\
0	NTM3	X1	A28	WHOLESALE	Large		
410215	NTM3	X1	A24	DIRECT	Small		
227368	NTM2	X1	A11	DIRECT	Large		
136731	NTM1	X1	A9	DIRECT	Large		
227357	NTM3	X1	A24	DIRECT	Small		
...	...	...	...	...	...		
420994	NTM1	X2	A43	DIRECT	Small		
67975	NTM1	X2	A44	DIRECT	Large		
420992	NTM1	X1	A3	DIRECT	Small		
210469	NTM1	X2	A21	DIRECT	Small		
338840	NTM2	X1	A5	DIRECT	Large		

	Product Type	Month of Sourcing	Sourcing Cost
0	Powder	2021-05-01	10.158

410215	Powder	2021-05-01	64.463
227368	Liquid	2021-05-01	151.696
136731	Powder	2021-05-01	146.982
227357	Powder	2021-05-01	73.149
...	...	...	...
420994	Powder	2020-07-01	157.094
67975	Liquid	2020-07-01	0.001
420992	Powder	2020-07-01	136.924
210469	Powder	2020-07-01	71.853
338840	Powder	2020-07-01	136.469

[550176 rows x 8 columns]

```
[11]: df_train.shape
```

```
[11]: (550176, 8)
```

```
[12]: df_test=pd.read_excel('/content/drive/MyDrive/AP moller/DS_ML Coding Challenge_
↳Dataset (1).xlsx',sheet_name='Test Dataset')
df_test
```

```
[12]:
```

	ProductType	Manufacturer	Area	Code	Sourcing	Channel	Product	Size \
0	NTM1	X1	A1		DIRECT		Small	
1	NTM1	X1	A10		DIRECT		Large	
2	NTM1	X1	A10		ECOM		Large	
3	NTM1	X1	A11		DIRECT		Large	
4	NTM1	X1	A2		DIRECT		Large	
..	...	...	...		...	...		
91	NTM3	X1	A44		DIRECT		Small	
92	NTM3	X1	A8		DIRECT		Large	
93	NTM3	X1	A8		DIRECT		Small	
94	NTM3	X2	A20		DIRECT		Large	
95	NTM3	X3	A22		RETAIL		Large	

	Product	Type	Month of Sourcing	Sourcing Cost
0	Powder		2021-06-21	103.68
1	Powder		2021-06-21	155.75
2	Powder		2021-06-21	143.02
3	Powder		2021-06-21	139.39
4	Powder		2021-06-21	169.42
..	...		...	...
91	Liquid		2021-06-21	89.57
92	Powder		2021-06-21	114.57
93	Powder		2021-06-21	111.26
94	Powder		2021-06-21	32.32
95	Powder		2021-06-21	40.73

[96 rows x 8 columns]

```
[13]: df_test = df_test.dropna(how='all')
df_test
```

```
[13]:   ProductType Manufacturer Area Code Sourcing Channel Product Size \
0          NTM1          X1      A1      DIRECT      Small
1          NTM1          X1     A10      DIRECT      Large
2          NTM1          X1     A10      ECOM      Large
3          NTM1          X1     A11      DIRECT      Large
4          NTM1          X1      A2      DIRECT      Large
..          ...          ...    ...    ...          ...
91         NTM3          X1     A44      DIRECT      Small
92         NTM3          X1      A8      DIRECT      Large
93         NTM3          X1      A8      DIRECT      Small
94         NTM3          X2     A20      DIRECT      Large
95         NTM3          X3     A22      RETAIL      Large
```

```
   Product Type Month of Sourcing Sourcing Cost
0      Powder  2021-06-21      103.68
1      Powder  2021-06-21      155.75
2      Powder  2021-06-21      143.02
3      Powder  2021-06-21      139.39
4      Powder  2021-06-21      169.42
..          ...          ...    ...
91     Liquid  2021-06-21      89.57
92     Powder  2021-06-21      114.57
93     Powder  2021-06-21      111.26
94     Powder  2021-06-21      32.32
95     Powder  2021-06-21      40.73
```

[96 rows x 8 columns]

```
[14]: TrainTestCombined = pd.concat([df_train, df_test], ignore_index=True)
TrainTestCombined
```

```
[14]:   ProductType Manufacturer Area Code Sourcing Channel Product Size \
0          NTM3          X1     A28  WHOLESALE      Large
1          NTM3          X1     A24      DIRECT      Small
2          NTM2          X1     A11      DIRECT      Large
3          NTM1          X1      A9      DIRECT      Large
4          NTM3          X1     A24      DIRECT      Small
...          ...          ...    ...    ...          ...
550267        NTM3          X1     A44      DIRECT      Small
550268        NTM3          X1      A8      DIRECT      Large
550269        NTM3          X1      A8      DIRECT      Small
550270        NTM3          X2     A20      DIRECT      Large
```

550271	NTM3	X3	A22	RETAIL	Large
--------	------	----	-----	--------	-------

	Product Type	Month of Sourcing	Sourcing Cost
0	Powder	2021-05-01	10.158
1	Powder	2021-05-01	64.463
2	Liquid	2021-05-01	151.696
3	Powder	2021-05-01	146.982
4	Powder	2021-05-01	73.149
...	...	...	...
550267	Liquid	2021-06-21	89.570
550268	Powder	2021-06-21	114.570
550269	Powder	2021-06-21	111.260
550270	Powder	2021-06-21	32.320
550271	Powder	2021-06-21	40.730

[550272 rows x 8 columns]

##number of duplicate rows

```
[15]: duplicate_rows_train = df_train[df_train.duplicated()]

# Get the number of duplicate rows
num_duplicate_rows_train = len(duplicate_rows_train)

print("Number of duplicate rows in training set:", num_duplicate_rows_train)
```

Number of duplicate rows in training set: 541165

```
[16]: duplicate_rows_test = df_test[df_test.duplicated()]

# Get the number of duplicate rows
num_duplicate_rows_test = len(duplicate_rows_test)

print("Number of duplicate rows in testing set:", num_duplicate_rows_test)
```

Number of duplicate rows in testing set: 0

```
[17]: duplicate_rows = TrainTestCombined[TrainTestCombined.duplicated()]

# Get the number of duplicate rows
num_duplicate_rows = len(duplicate_rows)

print("Number of duplicate rows in combined dataset:", num_duplicate_rows)
```

Number of duplicate rows in combined dataset: 541165

```
[18]: duplicate_rows
```

```
[18]:
```

	ProductType	Manufacturer	Area	Code	Sourcing	Channel	Product	Size	\
20	NTM1		X1	A1		DIRECT		Small	
31	NTM1		X1	A33		DIRECT		Large	
46	NTM2		X1	A6		DIRECT		Large	
52	NTM1		X2	A40		DIRECT		Large	
55	NTM2		X2	A31		DIRECT		Large	
...	...	...	...	...	...	...	...	...	...
550171	NTM1		X2	A43		DIRECT		Small	
550172	NTM1		X2	A44		DIRECT		Large	
550173	NTM1		X1	A3		DIRECT		Small	
550174	NTM1		X2	A21		DIRECT		Small	
550175	NTM2		X1	A5		DIRECT		Large	

	Product	Type	Month of Sourcing	Sourcing Cost
20		Powder	2021-05-01	114.899
31		Powder	2021-05-01	133.300
46		Powder	2021-05-01	144.391
52		Liquid	2021-05-01	24.479
55		Powder	2021-05-01	4.428
...	...	...	...	...
550171		Powder	2020-07-01	157.094
550172		Liquid	2020-07-01	0.001
550173		Powder	2020-07-01	136.924
550174		Powder	2020-07-01	71.853
550175		Powder	2020-07-01	136.469

[541165 rows x 8 columns]

```
[19]: (550176-541165)
```

```
[19]: 9011
```

```
[20]: (550272-542796)
```

```
[20]: 7476
```

```
[23]: df=df_train.copy()
```

```
##Number of null values
```

```
[24]: null_summary = df.isnull().sum()
null_summary
```

```
[24]: ProductType      0
Manufacturer      0
Area Code        0
Sourcing Channel  0
Product Size     0
```

```

Product Type      0
Month of Sourcing 0
Sourcing Cost     0
dtype: int64

```

```
[25]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 550176 entries, 0 to 338840
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ProductType           550176 non-null object
1   Manufacturer           550176 non-null object
2   Area Code             550176 non-null object
3   Sourcing Channel      550176 non-null object
4   Product Size          550176 non-null object
5   Product Type          550176 non-null object
6   Month of Sourcing     550176 non-null datetime64[ns]
7   Sourcing Cost         550176 non-null float64
dtypes: datetime64[ns](1), float64(1), object(6)
memory usage: 37.8+ MB

##describing the dataset

```

```
[26]: df.describe()
```

```

[26]:
count          Month of Sourcing  Sourcing Cost
mean  2020-12-08 10:27:28.769848576    108.816793
min    2020-07-01 00:00:00         -196.070000
25%    2020-10-01 00:00:00           57.000000
50%    2020-12-01 00:00:00          132.000000
75%    2021-03-01 00:00:00          146.147000
max    2021-05-01 00:00:00        32632.500000
std                  NaN           104.390097

```

```
##number of rows where the soourcing cost is zero
```

```
[27]: df[df["Sourcing Cost"]==0].count()
```

```

[27]: ProductType      10084
Manufacturer      10084
Area Code        10084
Sourcing Channel  10084
Product Size     10084
Product Type     10084
Month of Sourcing 10084
Sourcing Cost    10084

```

dtype: int64

##number of rows where the sourcing cost is negative

```
[28]: negative_entries = df[df['Sourcing Cost'] < 0]

print("Number of rows with negative sourcing cost:", len(negative_entries))
```

Number of rows with negative sourcing cost: 2231

##categorical columns

```
[29]: categorical_columns = df.select_dtypes(include=['object']).columns.tolist()

print("Categorical Columns:")
for col in categorical_columns:
    print(col)
```

Categorical Columns:

ProductType

Manufacturer

Area Code

Sourcing Channel

Product Size

Product Type

##unique entries in each categorical column

```
[30]: for col in categorical_columns:
        unique_entries = df[col].nunique()
        print(f"Number of unique entries in {col}: {unique_entries}")
```

Number of unique entries in ProductType: 3

Number of unique entries in Manufacturer: 3

Number of unique entries in Area Code: 45

Number of unique entries in Sourcing Channel: 4

Number of unique entries in Product Size: 3

Number of unique entries in Product Type: 2

##outlier detection

```
[31]: numeric_columns = df.select_dtypes(include=['number']).columns
numeric_columns
```

```
[31]: Index(['Sourcing Cost'], dtype='object')
```

##inter-quartile range

```
[32]: import pandas as pd
```



```
def detect_outliers_iqr(column):

    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    return (column < lower_bound) | (column > upper_bound)

numeric_columns = df.select_dtypes(include=['number']).columns
outliers_mask = df[numeric_columns].apply(detect_outliers_iqr)

outliers = df[outliers_mask.any(axis=1)]
print("Rows containing outliers:")
print(outliers)
```

Rows containing outliers:

	ProductType	Manufacturer	Area	Code	Sourcing Channel	Product Size	\
227129	NTM2	X1	A16		DIRECT	Small	
228770	NTM2	X1	A16		DIRECT	Small	
468492	NTM2	X1	A16		DIRECT	Small	
525026	NTM2	X1	A16		DIRECT	Small	
505209	NTM2	X1	A16		DIRECT	Small	
...	...	...	...	...	...	...	
214770	NTM2	X1	A23		RETAIL	Large	
62189	NTM2	X1	A37		DIRECT	Large	
63828	NTM2	X1	A37		DIRECT	Large	
422548	NTM1	X2	A42		DIRECT	Small	
70294	NTM2	X1	A37		DIRECT	Small	

	Product Type	Month of Sourcing	Sourcing Cost
227129	Powder	2021-05-01	1005.303
228770	Powder	2021-05-01	536.562
468492	Powder	2021-05-01	473.106
525026	Powder	2021-05-01	536.562
505209	Powder	2021-05-01	1005.303
...	...	...	...
214770	Powder	2020-07-01	2412.380
62189	Powder	2020-07-01	336.522
63828	Powder	2020-07-01	336.522
422548	Powder	2020-07-01	288.119
70294	Powder	2020-07-01	720.000

[2666 rows x 8 columns]

```
[33]: # Drop the rows containing outliers
cleaned_df = df[~outliers_mask.any(axis=1)]

# Display information about the removed outliers
print("Number of outliers removed:", outliers.shape[0])

# Display the cleaned DataFrame
print("DataFrame after removing outliers:")
print(cleaned_df)
```

Number of outliers removed: 2666

DataFrame after removing outliers:

	ProductType	Manufacturer	Area	Code	Sourcing	Channel	Product	Size	\
0	NTM3	X1	A28		WHOLESALE			Large	
410215	NTM3	X1	A24		DIRECT			Small	
227368	NTM2	X1	A11		DIRECT			Large	
136731	NTM1	X1	A9		DIRECT			Large	
227357	NTM3	X1	A24		DIRECT			Small	
...	...	...	...		...		...		
420994	NTM1	X2	A43		DIRECT			Small	
67975	NTM1	X2	A44		DIRECT			Large	
420992	NTM1	X1	A3		DIRECT			Small	
210469	NTM1	X2	A21		DIRECT			Small	
338840	NTM2	X1	A5		DIRECT			Large	

	Product	Type	Month of Sourcing	Sourcing Cost
0	Powder		2021-05-01	10.158
410215	Powder		2021-05-01	64.463
227368	Liquid		2021-05-01	151.696
136731	Powder		2021-05-01	146.982
227357	Powder		2021-05-01	73.149
...	...		...	...
420994	Powder		2020-07-01	157.094
67975	Liquid		2020-07-01	0.001
420992	Powder		2020-07-01	136.924
210469	Powder		2020-07-01	71.853
338840	Powder		2020-07-01	136.469

[547510 rows x 8 columns]

```
[34]: df=cleaned_df
len(df)
```

[34]: 547510

##Z-score method

```
[35]: from scipy import stats

# Function to detect outliers using Z-score
def detect_outliers_zscore(column, threshold=3):
    z_scores = stats.zscore(column)
    return abs(z_scores) > threshold

# Apply outlier detection to numeric columns in the DataFrame
outliers_mask = df[numeric_columns].apply(detect_outliers_zscore)
outliers = df[outliers_mask.any(axis=1)]
print("Rows containing outliers:")
print(len(outliers))
```

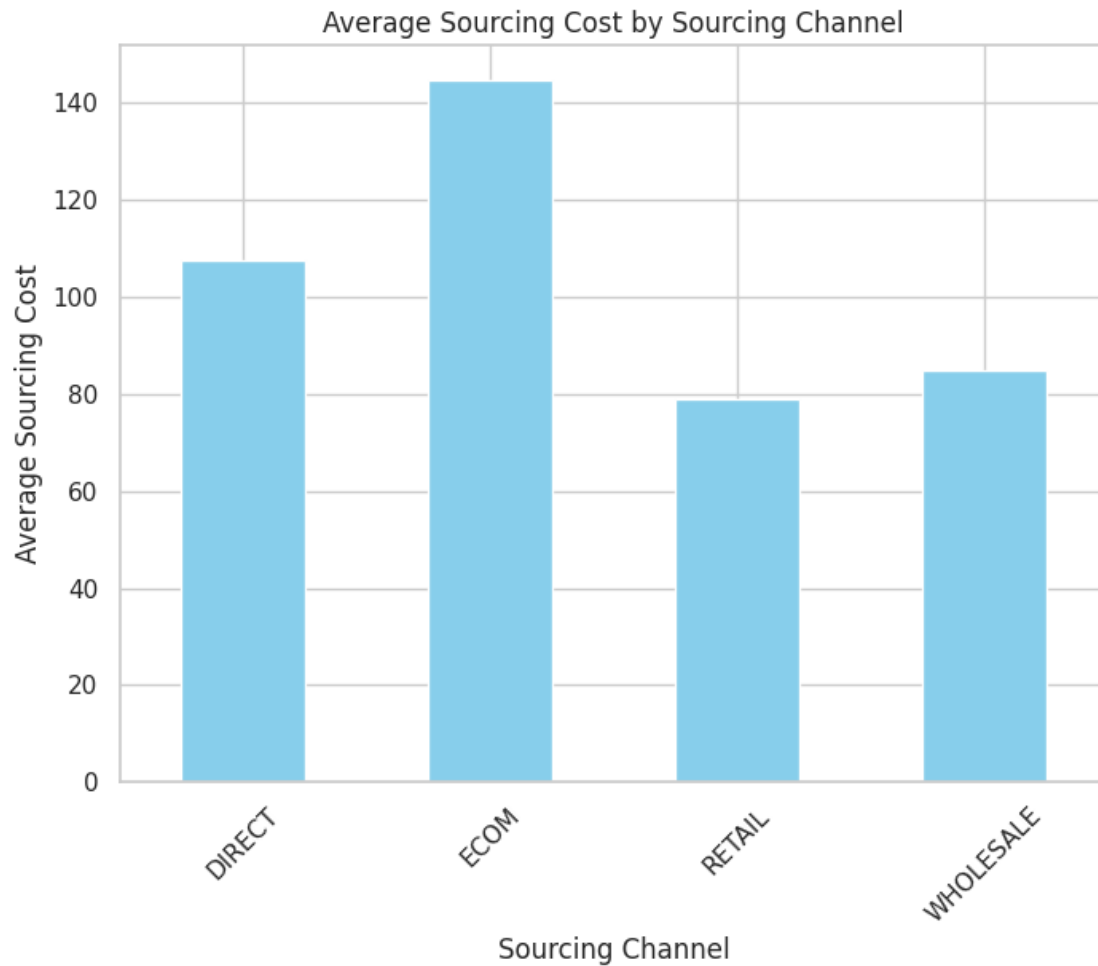
Rows containing outliers:

24

#Exploratory Data Analysis


```
[100]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
[107]: # Bar plot of average 'Sourcing Cost' by 'Sourcing Channel'
plt.figure(figsize=(8, 6))
df.groupby('Sourcing Channel')['Sourcing Cost'].mean().plot(kind='bar',
    color='skyblue')
plt.xlabel('Sourcing Channel')
plt.ylabel('Average Sourcing Cost')
plt.title('Average Sourcing Cost by Sourcing Channel')
plt.xticks(rotation=45)
plt.show()
```



The Sourcing Channel 'ECOM' tends to have the highest average Sourcing Cost compared to other channels in the dataset.

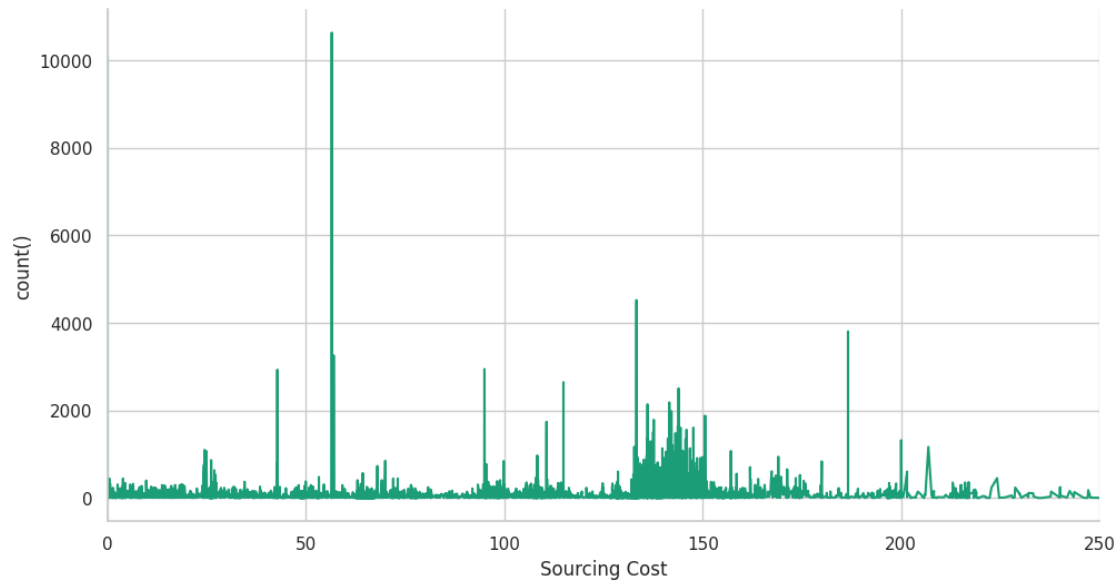
```
[108]: #Time series plot of 'Sourcing Cost' over time(month)  
#Note- here i had to downsample the cost value due to the number of values  
↳being too large to plot  
monthly_sourcing_cost = df.groupby('Month of Sourcing')['Sourcing Cost'].mean()  
plt.figure(figsize=(10, 6))  
plt.plot(monthly_sourcing_cost.index, monthly_sourcing_cost.values, marker='o',  
↳linestyle='-')  
plt.xlabel('Month of Sourcing')  
plt.ylabel('Average Sourcing Cost')  
plt.title('Average Sourcing Cost Over Time')  
plt.xticks(rotation=45)  
plt.grid(True)  
plt.show()
```




```
[109]: from matplotlib import pyplot as plt
import seaborn as sns
def _plot_series(series, series_name, series_index=0):
    palette = list(sns.palettes.mpl_palette('Dark2'))
    counted = (series['Sourcing Cost']
               .value_counts()
               .reset_index(name='counts')
               .rename({'index': 'Sourcing Cost'}, axis=1)
               .sort_values('Sourcing Cost', ascending=True))
    xs = counted['Sourcing Cost']
    ys = counted['counts']
    plt.plot(xs, ys, label=series_name, color=palette[series_index % len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = df.sort_values('Sourcing Cost', ascending=True)
_plot_series(df_sorted, '')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Sourcing Cost')
_ = plt.ylabel('count()')
plt.xlim(0,250)
```

[109]: (0.0, 250.0)

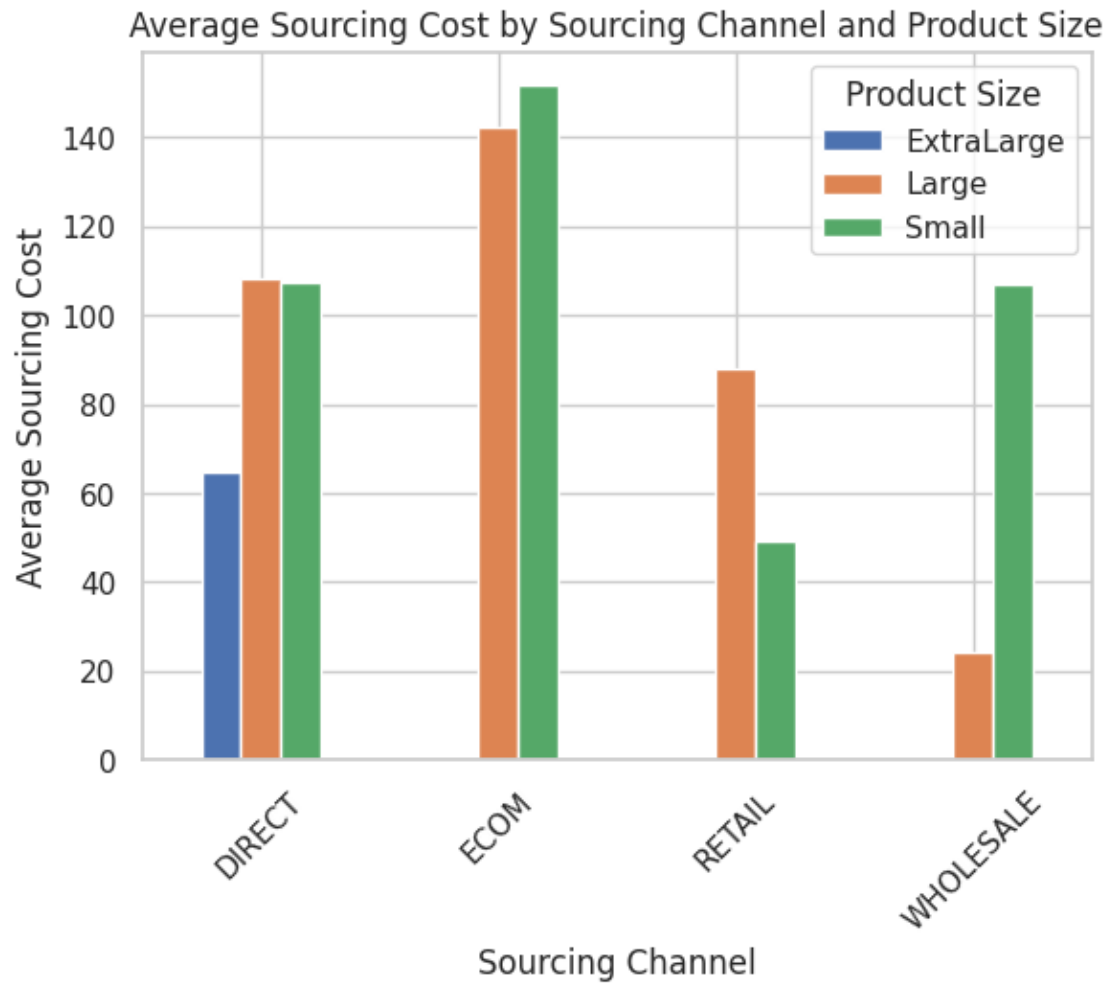


```
[110]: # Pie chart of manufacturers
plt.figure(figsize=(8, 8))
manufacturer_counts = df['Manufacturer'].value_counts()
plt.pie(manufacturer_counts, labels=manufacturer_counts.index, autopct='%1.
↪1f%%', startangle=140)
plt.title('Distribution of Manufacturers')
plt.axis('equal')
plt.show()
```



```
[114]: custom_labels = {1: 'Large', 2: 'Small', 0: 'Extra Large'}
df.groupby(['Sourcing Channel', 'Product Size'])['Sourcing Cost'].mean().
    ↪unstack().plot(kind='bar')

plt.xlabel('Sourcing Channel')
plt.ylabel('Average Sourcing Cost')
plt.title('Average Sourcing Cost by Sourcing Channel and Product Size')
plt.xticks(rotation=45)
plt.show()
```



##Preprocessing

- encoding
- feature engineering
- handle duplicate values
- scaling of numerical features (when passing to model)
- missing values(no missing values)
- handling outliers
- negative entries in sourcing cost

[44]: `df.head()`

```
[44]:      ProductType Manufacturer Area Code Sourcing Channel Product Size \
0          NTM3          X1      A28      WHOLESALE      Large
```



410215	NTM3	X1	A24	DIRECT	Small
227368	NTM2	X1	A11	DIRECT	Large
136731	NTM1	X1	A9	DIRECT	Large
227357	NTM3	X1	A24	DIRECT	Small

	Product Type	Month of Sourcing	Sourcing Cost
0	Powder	2021-05-01	10.158
410215	Powder	2021-05-01	64.463
227368	Liquid	2021-05-01	151.696
136731	Powder	2021-05-01	146.982
227357	Powder	2021-05-01	73.149

```
[45]: df_training = df.copy()
```

```
[46]: df_testing = df_test.copy()
```

```
[47]: from sklearn.preprocessing import LabelEncoder
def label_encode_categorical_columns(df, categorical_columns):
    # Make a copy of the DataFrame to avoid modifying the original
    df_encoded = df.copy()

    # Initialize LabelEncoder
    label_encoder = LabelEncoder()

    # Iterate over each categorical column and perform label encoding
    for col in categorical_columns:
        # Fit label encoder and transform the column
        df_encoded[col] = label_encoder.fit_transform(df_encoded[col])

    return df_encoded
```

```
[48]: from sklearn.preprocessing import MinMaxScaler
def Normalize_column(df, numeric_columns):
    scaler = MinMaxScaler()
    # Fit and transform the scaler on the numerical columns
    df[numeric_columns] = scaler.fit_transform(df[numeric_columns])
    return df
```

```
[63]: label_encoder = LabelEncoder()
df_training = label_encode_categorical_columns(df_training, categorical_columns)

df_training['Month of Sourcing'] = label_encoder.
    ↪ fit_transform(df_training['Month of Sourcing'])
df_testing = label_encode_categorical_columns(df_testing, categorical_columns)
df_testing['Month of Sourcing'] = label_encoder.fit_transform(df_testing['Month_
    ↪ of Sourcing'])
```

```
[64]: df_training.head()
```

```
[64]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size \
0	2	0	19	3	1
410215	2	0	16	0	2
227368	1	0	2	0	1
136731	0	0	44	0	1
227357	2	0	16	0	2

	Product Type	Month of Sourcing	Sourcing Cost
0	1	10	0.200889
410215	1	10	0.361896
227368	0	10	0.620529
136731	1	10	0.606552
227357	1	10	0.387648

```
[65]: df_testing.head()
```

```
[65]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size \
0	0	0	0	0	2
1	0	0	1	0	1
2	0	0	1	1	1
3	0	0	2	0	1
4	0	0	11	0	1

	Product Type	Month of Sourcing	Sourcing Cost
0	1	0	0.431713
1	1	0	0.657544
2	1	0	0.602333
3	1	0	0.586590
4	1	0	0.716832

```
[66]: df_training= Normalize_column(df_training, numeric_columns)
df_testing= Normalize_column(df_testing, numeric_columns)
```

```
[67]: df_training.head()
```

```
[67]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size \
0	2	0	19	3	1
410215	2	0	16	0	2
227368	1	0	2	0	1
136731	0	0	44	0	1
227357	2	0	16	0	2

	Product Type	Month of Sourcing	Sourcing Cost
0	1	10	0.200889
410215	1	10	0.361896

227368	0	10	0.620529
136731	1	10	0.606552
227357	1	10	0.387648

```
[68]: df_testing.head()
```

```
[68]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size \
0	0	0	0	0	2
1	0	0	1	0	1
2	0	0	1	1	1
3	0	0	2	0	1
4	0	0	11	0	1

	Product Type	Month of Sourcing	Sourcing Cost
0	1	0	0.431713
1	1	0	0.657544
2	1	0	0.602333
3	1	0	0.586590
4	1	0	0.716832

##traditional machine learning

```
[69]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
import matplotlib.pyplot as plt

# Assuming your preprocessed dataset is stored in a pandas DataFrame named 'df'

# Split the dataset into training and testing sets
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df_training.drop(["Sourcing_
↪Cost"], axis=1), df_training["Sourcing Cost"], test_size=0.2,
↪random_state=42)
```

```
[70]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
import matplotlib.pyplot as plt
```

```

# Assuming your preprocessed dataset is stored in a pandas DataFrame named 'df'

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df_training.drop(["Sourcing_
↪Cost"],axis=1), df_training["Sourcing Cost"], test_size=0.2, random_state=42)

```

```

[71]: # Initialize models
linear_reg = LinearRegression()
decision_tree = DecisionTreeRegressor()
random_forest = RandomForestRegressor()
xgboost = XGBRegressor()

# Fit models
linear_reg.fit(X_train, y_train)
decision_tree.fit(X_train, y_train)
random_forest.fit(X_train, y_train)
xgboost.fit(X_train, y_train)

# Make predictions
y_pred_linear_reg = linear_reg.predict(X_test)
y_pred_decision_tree = decision_tree.predict(X_test)
y_pred_random_forest = random_forest.predict(X_test)
y_pred_xgboost = xgboost.predict(X_test)

# Calculate error metrics
def calculate_error_metrics(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = mean_squared_error(y_true, y_pred, squared=False)
    r2 = r2_score(y_true, y_pred)
    return mae, mse, rmse, r2

mae_linear_reg, mse_linear_reg, rmse_linear_reg, r2_linear_reg = ↪
↪calculate_error_metrics(y_test, y_pred_linear_reg)
mae_decision_tree, mse_decision_tree, rmse_decision_tree, r2_decision_tree = ↪
↪calculate_error_metrics(y_test, y_pred_decision_tree)
mae_random_forest, mse_random_forest, rmse_random_forest, r2_random_forest = ↪
↪calculate_error_metrics(y_test, y_pred_random_forest)
mae_xgboost, mse_xgboost, rmse_xgboost, r2_xgboost = ↪
↪calculate_error_metrics(y_test, y_pred_xgboost)

# Plot evaluation metrics
models = ['Linear Regression', 'Decision Tree', 'Random Forest', 'XGBoost']
mae_values = [mae_linear_reg, mae_decision_tree, mae_random_forest, mae_xgboost]
mse_values = [mse_linear_reg, mse_decision_tree, mse_random_forest, mse_xgboost]
rmse_values = [rmse_linear_reg, rmse_decision_tree, rmse_random_forest, ↪
↪rmse_xgboost]

```

```

r2_values = [r2_linear_reg, r2_decision_tree, r2_random_forest, r2_xgboost]

plt.figure(figsize=(10, 6))

plt.subplot(2, 2, 1)
plt.bar(models, mae_values)
plt.title('Mean Absolute Error')
plt.ylabel('MAE')

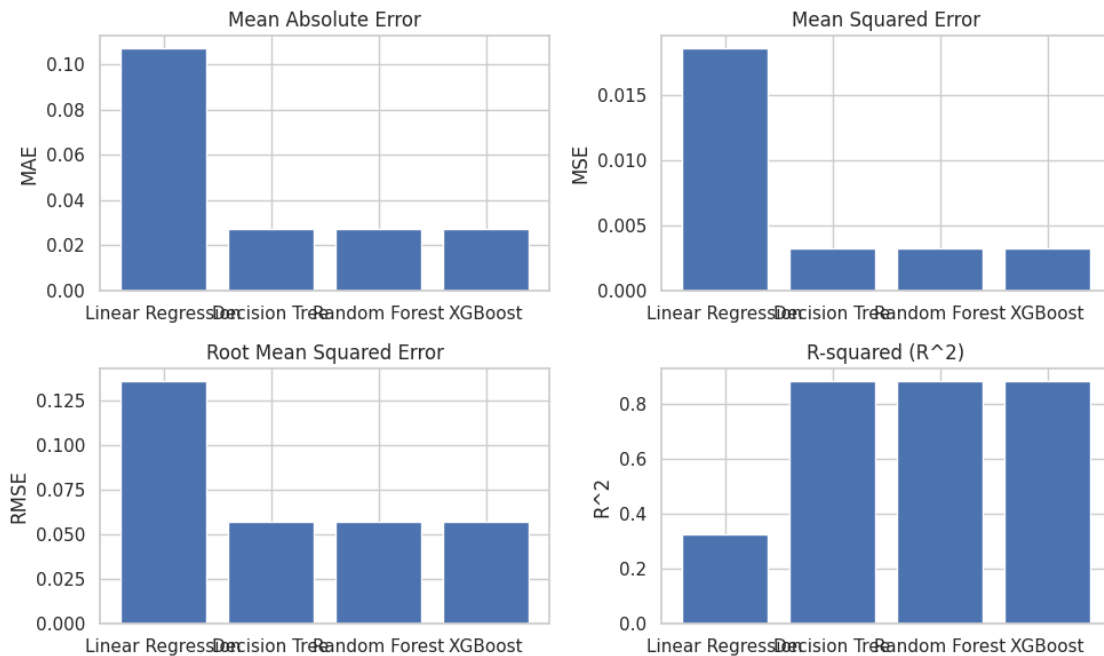
plt.subplot(2, 2, 2)
plt.bar(models, mse_values)
plt.title('Mean Squared Error')
plt.ylabel('MSE')

plt.subplot(2, 2, 3)
plt.bar(models, rmse_values)
plt.title('Root Mean Squared Error')
plt.ylabel('RMSE')

plt.subplot(2, 2, 4)
plt.bar(models, r2_values)
plt.title('R-squared (R^2)')
plt.ylabel('R^2')

plt.tight_layout()
plt.show()

```



```
[72]: r2_values
```

```
[72]: [0.3231536406689248,  
      0.8821268067235862,  
      0.8821184862783993,  
      0.8815434326345656]
```

```
##Testing
```

```
[73]: df_testing_1=df_testing.drop(["Sourcing Cost"],axis=1)  
      df_testing_1["Month of Sourcing"]=11  
      df_testing_1.head()
```

```
[73]:
```

	ProductType	Manufacturer	Area	Code	Sourcing Channel	Product Size	\
0	0	0		0	0	2	
1	0	0		1	0	1	
2	0	0		1	1	1	
3	0	0		2	0	1	
4	0	0		11	0	1	

	Product Type	Month of Sourcing
0	1	11
1	1	11
2	1	11
3	1	11
4	1	11

```
[74]: y_test_pred_random_forest = random_forest.predict(df_testing_1)
```

```
[75]: y_true=df_testing["Sourcing Cost"]  
      y_true
```

```
[75]: 0      0.431713  
      1      0.657544  
      2      0.602333  
      3      0.586590  
      4      0.716832  
      ...  
      91     0.370517  
      92     0.478943  
      93     0.464588  
      94     0.122219  
      95     0.158694  
      Name: Sourcing Cost, Length: 96, dtype: float64
```

```
[76]: r2 = r2_score(y_true, y_test_pred_random_forest)
r2
```

```
[76]: 0.6088622798782557
```

```
##Deep learning approaches for time series
```

```
[77]: df_training.drop(columns=['Sourcing Cost'])
```

```
[77]:
```

	ProductType	Manufacturer	Area	Code	Sourcing Channel	Product Size	\
0	2	0	19		3	1	
410215	2	0	16		0	2	
227368	1	0	2		0	1	
136731	0	0	44		0	1	
227357	2	0	16		0	2	
...	...	...	...		...	...	
420994	0	1	36		0	2	
67975	0	1	37		0	1	
420992	0	0	21		0	2	
210469	0	1	13		0	2	
338840	1	0	40		0	1	

	Product Type	Month of Sourcing
0	1	10
410215	1	10
227368	0	10
136731	1	10
227357	1	10
...	...	...
420994	1	0
67975	0	0
420992	1	0
210469	1	0
338840	1	0

```
[547510 rows x 7 columns]
```

```
#LSTM(Best model)
```

```
[78]: # Step 1: Import necessary libraries
import pandas as pd
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
```

```

# Step 3: Prepare data
X = df_training.drop(columns=['Sourcing Cost']).values # Input features (all
↳ columns except 'Sourcing Cost')
y = df_training['Sourcing Cost'].values # Target variable

# Reshape input features for LSTM (assuming a single time step)
# The reshape is necessary for LSTM input (samples, time steps, features)
X = X.reshape(X.shape[0], 1, X.shape[1])

# Step 4: Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Step 8: Make forecasts
# You can make forecasts using model.predict() on new data
# Example: y_pred = model.predict(X_new_data)

```

```
[79]: X_train.shape
```

```
[79]: (438008, 1, 7)
```

```

[80]: lstm_model = Sequential([
        LSTM(units=50, input_shape=(X_train.shape[1], X_train.shape[2])),
        Dense(1)
    ])

lstm_model.compile(optimizer='adam', loss='mse')
lstm_model.fit(X_train, y_train, epochs=25, batch_size=32, verbose=1)

loss = lstm_model.evaluate(X_test, y_test, verbose=0)
print(f'Test Loss: {loss}')

```

```

Epoch 1/25
13688/13688 [=====] - 52s 3ms/step - loss: 0.0104
Epoch 2/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0083
Epoch 3/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0079
Epoch 4/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0075
Epoch 5/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0073
Epoch 6/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0070
Epoch 7/25
13688/13688 [=====] - 46s 3ms/step - loss: 0.0067

```



```

Epoch 8/25
13688/13688 [=====] - 47s 3ms/step - loss: 0.0063
Epoch 9/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0060
Epoch 10/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0056
Epoch 11/25
13688/13688 [=====] - 46s 3ms/step - loss: 0.0053
Epoch 12/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0051
Epoch 13/25
13688/13688 [=====] - 47s 3ms/step - loss: 0.0050
Epoch 14/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0049
Epoch 15/25
13688/13688 [=====] - 45s 3ms/step - loss: 0.0048
Epoch 16/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0048
Epoch 17/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0047
Epoch 18/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0046
Epoch 19/25
13688/13688 [=====] - 46s 3ms/step - loss: 0.0045
Epoch 20/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0045
Epoch 21/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0044
Epoch 22/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0044
Epoch 23/25
13688/13688 [=====] - 44s 3ms/step - loss: 0.0043
Epoch 24/25
13688/13688 [=====] - 43s 3ms/step - loss: 0.0043
Epoch 25/25
13688/13688 [=====] - 46s 3ms/step - loss: 0.0043
Test Loss: 0.004411341156810522

```

```

[ ]: model_file_path = 'lstm_model.h5'

# Save the model
lstm_model.save(model_file_path)

print("LSTM model saved successfully at:", model_file_path)

```

```

[81]: import matplotlib.pyplot as plt

```

```


# Step 1: Make predictions using the trained LSTM model
y_pred = lstm_model.predict(X_test)

# Step 2: Extract the first 100 data points from the original dataset and LSTM_
↳ predictions
original_values = y_test[:100] # Original values
predicted_values = y_pred[:100] # LSTM predictions

# Step 3: Plot the original values and LSTM predictions
plt.figure(figsize=(10, 6))
plt.plot(original_values, label='Original Values', marker='o')
plt.plot(predicted_values, label='LSTM Predictions', marker='x')
plt.title('Comparison of Original Values and LSTM Predictions')
plt.xlabel('Data Point Index')
plt.ylabel('Sourcing Cost')
plt.legend()
plt.grid(True)
plt.show()

```

3422/3422 [=====] - 8s 2ms/step



```

[82]: r2 = r2_score(original_values, predicted_values)
      r2

```

```
[82]: 0.8421635064329565
```

```
###TESTING DATASET
```

```
[88]: from sklearn.preprocessing import LabelEncoder

# Assuming your training set DataFrame is named 'train_df'
# Assuming 'categorical_columns' is a list containing the names of categorical
↳ columns

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Iterate over each categorical column and perform label encoding
for col in categorical_columns:
    # Fit label encoder and transform the column
    df_test[col] = label_encoder.fit_transform(df_test[col])
```

```
[89]: df_test.head()
```

```
[89]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size	\
0	0	0	0	0	2	
1	0	0	1	0	1	
2	0	0	1	1	1	
3	0	0	2	0	1	
4	0	0	11	0	1	

	Product Type	Month of Sourcing	Sourcing Cost
0	1	2021-06-21	103.68
1	1	2021-06-21	155.75
2	1	2021-06-21	143.02
3	1	2021-06-21	139.39
4	1	2021-06-21	169.42

```
[90]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the scaler on the numerical columns
df_test[numeric_columns] = scaler.fit_transform(df_test[numeric_columns])
```

```
[91]: df_test
```

```
[91]:
```

	ProductType	Manufacturer	Area Code	Sourcing Channel	Product Size	\
0	0	0	0	0	2	
1	0	0	1	0	1	
2	0	0	1	1	1	
3	0	0	2	0	1	
4	0	0	11	0	1	

```

..          ...          ...          ...          ...          ...
91          2          0          37          0          2
92          2          0          43          0          1
93          2          0          43          0          2
94          2          1          12          0          1
95          2          2          14          2          1

```

```

      Product Type Month of Sourcing Sourcing Cost
0          1      2021-06-21      0.431713
1          1      2021-06-21      0.657544
2          1      2021-06-21      0.602333
3          1      2021-06-21      0.586590
4          1      2021-06-21      0.716832
..          ...          ...          ...
91          0      2021-06-21      0.370517
92          1      2021-06-21      0.478943
93          1      2021-06-21      0.464588
94          1      2021-06-21      0.122219
95          1      2021-06-21      0.158694

```

[96 rows x 8 columns]

```
[92]: df_test["Month of Sourcing"] = label_encoder.fit_transform(df_test["Month of_
↪Sourcing"])
```

```
[93]: df_test.head()
```

```
[93]:   ProductType  Manufacturer  Area Code  Sourcing Channel  Product Size \
0          0          0          0          0          2
1          0          0          1          0          1
2          0          0          1          1          1
3          0          0          2          0          1
4          0          0          11         0          1

```

```

      Product Type Month of Sourcing Sourcing Cost
0          1      2021-06-21      0.431713
1          1      2021-06-21      0.657544
2          1      2021-06-21      0.602333
3          1      2021-06-21      0.586590
4          1      2021-06-21      0.716832

```

```
[94]: X = df_test.drop(columns=['Sourcing Cost']).values # Input features (all_
↪columns except 'Sourcing Cost')
y = df_test['Sourcing Cost'].values # Target variable

# Reshape input features for LSTM (assuming a single time step)
# The reshape is necessary for LSTM input (samples, time steps, features)
```

```
X = X.reshape(X.shape[0], 1, X.shape[1])
```

```
[95]: X.shape
```

```
[95]: (96, 1, 7)
```


```
[96]: import matplotlib.pyplot as plt

# Step 1: Make predictions using the trained LSTM model
y_pred = lstm_model.predict(X)

# Step 2: Extract the first 100 data points from the original dataset and LSTM
↳ predictions
original_values = y[:-1] # Original values
predicted_values = y_pred[:-1] # LSTM predictions

# Step 3: Plot the original values and LSTM predictions
plt.figure(figsize=(10, 6))
plt.plot(original_values, label='Original Values', marker='o')
plt.plot(predicted_values, label='LSTM Predictions', marker='x')
plt.title('Comparison of Original Values and LSTM Predictions')
plt.xlabel('Data Point Index')
plt.ylabel('Sourcing Cost')
plt.legend()
plt.grid(True)
plt.show()
```

```
3/3 [=====] - 0s 4ms/step
```



as we can see, LSTM is able to classify and forecast the points quite accurately

[ ]: